

# Data Engineer Take Home Test 2026

---

## Introduction

In this test, you will design and implement an end-to-end data pipeline from source to destination. We have set up **PostgreSQL** and **Kafka** for you to use in this assignment.

Your main objectives are:

1. Use PostgreSQL as both source and data warehouse
2. Design a **batch pipeline** for daily data processing
3. Design a **streaming pipeline** for real-time data processing
4. Demonstrate your understanding of ETL principles and containerization

**Assessment Criteria:** We will evaluate your understanding of Data Engineering fundamentals, code quality, and ability to work with modern data tools.

---


## Files Included

```
Shell
assignment/
├── docker-compose.yaml      # Container orchestration
├── kafka.py                 # Script to load data to Kafka
├── pg.py                    # Script to load data to PostgreSQL
├── requirements.txt          # Python dependencies
├── products.csv              # Product master data
├── orders.csv                # Order transaction data
├── problem-0/
│   ├── products.sql           # (Fill this) Products table
│   └── orders.sql             # (Fill this) Orders table schema
└── problem-1/                # Batch pipeline folder
    └── answer.md              # (Fill this) SQL answers
└── problem-2/                # Streaming pipeline folder
```

# Data Overview

## Architecture

Source (PostgreSQL/Kafka) → Pipeline (Your Code) → Destination (PostgreSQL)



## Products Data (products.csv)

Sample product catalog:

product_id	product_name	price	category
PROD-001	Coffee	45	Beverage
PROD-002	Green Tea	40	Beverage
PROD-003	Sandwich	85	Food

## Orders Data (orders.csv)

Sample order transactions:

order_id	order_timestamp	user_id	product_id	quantity	status
ORD-00001	2024-01-01 11:00:00	USER-001	PROD-001	2	COMPLETE
ORD-00002	2024-01-01 12:00:00	USER-002	PROD-003	1	COMPLETE
ORD-00003	2024-01-02 08:00:00	USER-001	PROD-002	1	COMPLETE

**Note:** Only orders with status = 'COMPLETE' should be included in revenue calculations.

## Problem 0 - Environment Setup

---

### Tasks

1. **Install Docker** on your local machine (if not already installed)
  2. **Start services** using docker-compose up -d
  3. **Verify Kafka UI** is accessible at <http://localhost:8080>
  4. **Define table schemas** in the following files:
    - a. problem-0/products.sql
    - b. problem-0/orders.sql
  5. **Load data to Kafka** by running python kafka.py
  6. **Load data to PostgreSQL** by running python pg.py
  7. **Verify data** is correctly loaded in both Kafka and PostgreSQL
- 

### Kafka UI Verification

After running kafka.py, you should see:

- Topic: orders (with messages)
  - Topic: products (with messages)
- 

### PostgreSQL Verification

After running pg.py, verify with:

```
SELECT COUNT(*) FROM orders; – Result: 1000000 records  
SELECT COUNT(*) FROM products; – Result: 10 records  
SELECT * FROM orders LIMIT 5;  
SELECT * FROM products LIMIT 5;
```

---

### Deliverables

Two **SQL files** with CREATE TABLE statements:

- [assignment/problem-0/sql/products.sql](#)
- [assignment/problem-0/sql/orders.sql](#)

### Hints:

- Look at the CSV files to understand column types
- Consider appropriate data types (VARCHAR, INTEGER, DATE, etc.)
- Think about which columns should be primary keys

# Problem 1 - Batch Pipeline

---

## Business Requirements

Your analytics team needs daily reports to monitor business performance. Design a batch ETL pipeline with these requirements:

1. **Schedule:** Data must be ready for querying every morning at 8 AM
  2. **Key Metrics to Calculate:**
    - a. Daily total revenue (only COMPLETE orders)
    - b. Daily number of new customers (first-time buyers)
    - c. Daily order count by product category
  3. **Incremental Processing:** Pipeline should process only new data each day (not reprocess everything)
- 

## Technical Requirements

- **Source:** Read from PostgreSQL tables (orders, products)
  - **Processing:** Apache Spark is mandatory for batch processing. Use PySpark to build your batch pipeline.
  - **Destination:** Write aggregated results to PostgreSQL (you design the schema)
  - **Orchestration:** Airflow (should include in docker)
- 

## SQL Questions

**Question 1:** Find the total revenue for user USER-001 across all completed orders.

**Question 2:** Find how many completed orders were placed on 2024-01-03.

**Question 3:** Find the average daily order count for the 'Beverage' product category in January 2024.

**Question 4:** Find the date with the highest number of new customers acquired.

## **Deliverables**

The deliverables are broken down into their respective directories:

### **1. Batch Pipeline Code (src/ and dags/):**

- **Core PySpark application logic**, including data extraction, transformation (joins, aggregations), loading to destination table(s), and incremental processing logic, placed under [assignment/problem-1/src/](#).
- **Airflow DAGs** for orchestration placed under [assignment/problem-1/dags/](#).
- **Containerization**: [assignment/problem-1/Dockerfile](#) and updated [assignment/problem-1/docker-compose.yml](#) for the pipeline.
- **Documentation**: [assignment/problem-1/README.md](#) explaining how to run the code.

### **2. Data Model (doc/ and sql/):**

- **Design Document**: [assignment/problem-1/doc/data\\_model.png](#) or [data\\_model.md](#) showing your destination table schema(s), explanation of design choices, and an ER diagram or data model diagram.
- **Destination Table DDL**: [assignment/problem-1/sql/problem-1-<table\\_name>.sql](#) (DDL for your new table(s)).
- **Solution Explanation**: [assignment/problem-1/doc/solution.md](#) explaining the overall batch solution.

### **3. SQL Queries + Answers (sql/ and results/):**

- **SQL Query Files**: Dedicated SQL files for the required questions, placed under [assignment/problem-1/sql/](#):
  - problem-1-ddl-<table\_name> ()
  - problem-1-dml-1.sql (Question 1)
  - problem-1-dml-2.sql (Question 2)
  - problem-1-dml-3.sql (Question 3)
  - problem-1-dml-4.sql (Question 4)
- **Results**: Screen captures of the SQL query results placed under [assignment/problem-1/results/](#):
  - problem-1-question-1.png
  - problem-1-question-2.png
    - ...and so on for all questions.

## Problem 2 - Streaming Pipeline

---

### Business Requirements

The business team wants **real-time monitoring** of order activity. They need to see:

- **Order count per minute** for each product
- **Update frequency:** Within 1 minute of order placement

This will power a real-time dashboard showing live order activity.

---

### Technical Requirements

- **Source:** Kafka topic orders
  - **Processing:** Streaming framework (e.g., Flink, or Spark Streaming)
  - **Output:** Can be to console output, or file
  - **Aggregation Window:** 1-minute tumbling window
- 

### Example Output

product_id	window_start	window_end	order_count
PROD-001	2024-01-01 10:00:00	2024-01-01 10:01:00	5
PROD-002	2024-01-01 10:00:00	2024-01-01 10:01:00	3
PROD-001	2024-01-01 10:01:00	2024-01-01 10:02:00	2

## **Deliverables**

The deliverables are organized under the problem-2/ directory:

### **1. Streaming Pipeline Code:**

- [assignment/problem-2/streaming\\_pipeline.py](#) (or your preferred language) containing the Kafka consumer setup, window aggregation logic, and output mechanism.

### **2. Documentation and Results:**

- **Solution Explanation:** [assignment/problem-2/doc/solution.md](#) explaining your streaming solution.
- **Results:** Screen capture of the example output in [assignment/problem-2/results/problem-2-question-1.png](#).
- **Documentation:** [assignment/problem-2/README.md](#) explaining how to run the streaming pipeline code.

## **Hints**

- Start with a simple Kafka consumer that prints messages
- Add windowing logic incrementally
- Test with the existing data first, then consider how it would work with real-time data
- You can simulate real-time data by producing messages slowly to Kafka

# Submission Guidelines

## File Structure

```
Shell
assignment/
└── problem-0/
    └── sql
        ├── products.sql
        └── orders.sql
    └── problem-1/
        ├── results # screen capture
        │   ├── problem-1-question-1.png # result for question 1
        │   ├── problem-1-question-2.png # result for question 2
        │   └── ...
        ├── dags/
        ├── plugins/
        ├── src/      # application code
        ├── tests/
        ├── Dockerfile
        ├── docker-compose.yml
        ├── doc/
        │   ├── solution.md # explain your solution
        │   └── data_model.png # data model image (or .md)
        ├── sql      # required SQL files
        │   ├── problem-1-ddl-<table_name>.sql    # ddl for your new table
        │   ├── problem-1-dml-1.sql    # sql for question 1
        │   ├── problem-1-dml-2.sql    # sql for question 2
        │   └── ...sql
        ├── config/
        └── README.md # how to run your code
    └── problem-2/
        ├── results # screen capture
        │   └── problem-2-question-1.png
        ├── doc/
        │   └── solution.md # explain your solution
        ├── streaming_pipeline.py # your streaming pipeline code
        └── README.md # how to run your code
└── (all original files)
```

## Submission Method

1. Ensure all files are in the assignment/ folder
2. Compress the folder: [firstname]-[lastname]-lmwn-de.zip
3. Submit via the provided link/email

**Example:** john-smith-lmwn-de.zip

## **Questions?**

If you have questions about the assignment requirements or encounter technical issues with the provided setup, please contact us.

**Good luck!**